

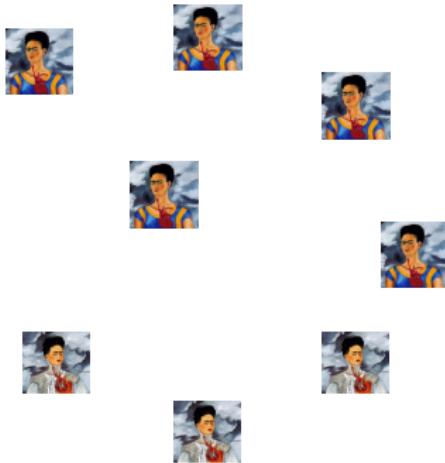
# Parameterized verification of asynchronous round-based distributed algorithms reduced to nuXmv

Nathalie Bertrand

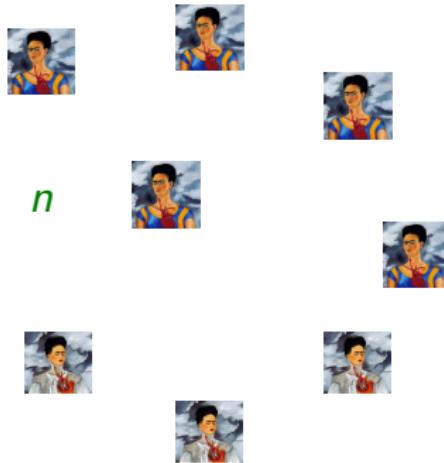


joint work with Pranav Ghorpade and Sasha Rubin  
University of Sydney

# Fault-tolerant distributed algorithms

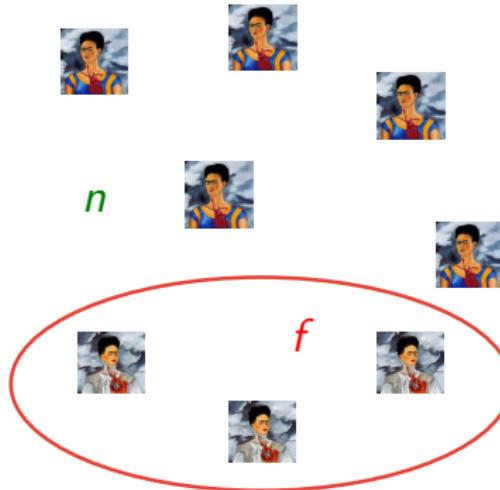


# Fault-tolerant distributed algorithms



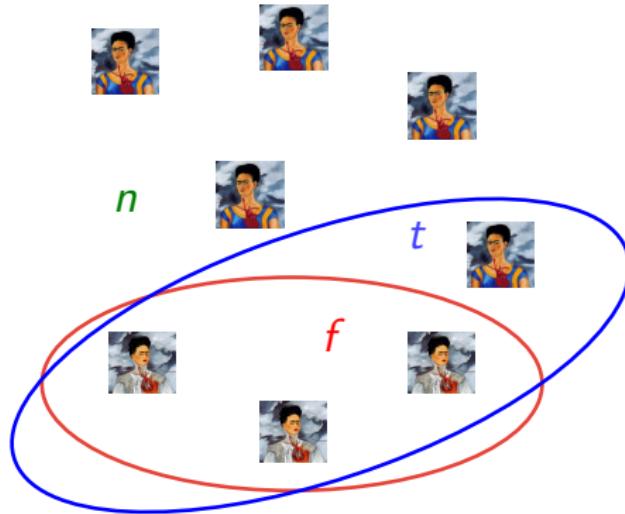
- $n$  processes

# Fault-tolerant distributed algorithms



- $n$  processes
- $f$  are faulty (e.g. crash or Byzantine failures)

# Fault-tolerant distributed algorithms



- $n$  processes
- $f$  are faulty (e.g. crash or Byzantine failures)
- $t$  known upper bound on  $f$
- resilience condition between these parameters, e.g.  $2t < n$

# Asynchronous round-based distributed algorithms

Consensus or leader election protocols

- asynchronous communication by broadcast
- threshold guards on number of received messages
- finitely many local variables
- structured in rounds:
  - rounds are identical up to round index, used to tag messages
  - round increment not limited to  $r := r + 1$

```
bool v := input_value({0,1});  
int r := 1;  
while (true) do  
    broadcast (v,r);  
    wait for n - t messages (*,r);  
    if received 2(n + t)/3 messages (w,r)  
    then d := w; halt  
    else if received (n + t)/2 messages (0,r)  
    then v := 0; r:=r+2;  
    else if received (n + t)/3 messages (1,r)  
    then v := 1; r:=r+1;  
od
```

# Existing formal methods approaches

No fully automated techniques; Mostly human-guided methods

- interactive theorem provers
  - TLA+ protocol formalization and verification for Paxos [Lamport, Merz, Doligez 2012], multi-Paxos [Chand, Liu, Stoller 2016] and DAG-based consensus in TLA+ [Bertrand, Ghorpade, Rubin, Scholz, Subotić 2025]
  - Rocq/VERDI specification and verification of Raft [Woos, Wilcox, Anton, Tatlock, Ernst, Anderson 2016]
- reduction to existing tools
  - restricted schedulers for randomized algorithms [Bertrand, Konnov, Lazić, Widder 2020]
- model checking with fixed number of processes
  - reduction theorem for finite instances to TLC [Chaouch-Saad, Charron-Bost, Merz 2009]
  - Paxos in SPIN [Delzanno, Tatarek, Traverso 2014]
  - agreement for asynchronous consensus algorithms [Noguchi, Tsuchiya, Kikuno 2012]

# Our approach

## Challenges

- 2 sources of infinity: number of processes, number of rounds
- asynchronous communications: unbounded *drift* between processes

# Our approach

## Challenges

- 2 sources of infinity: number of processes, number of rounds
- asynchronous communications: unbounded *drift* between processes



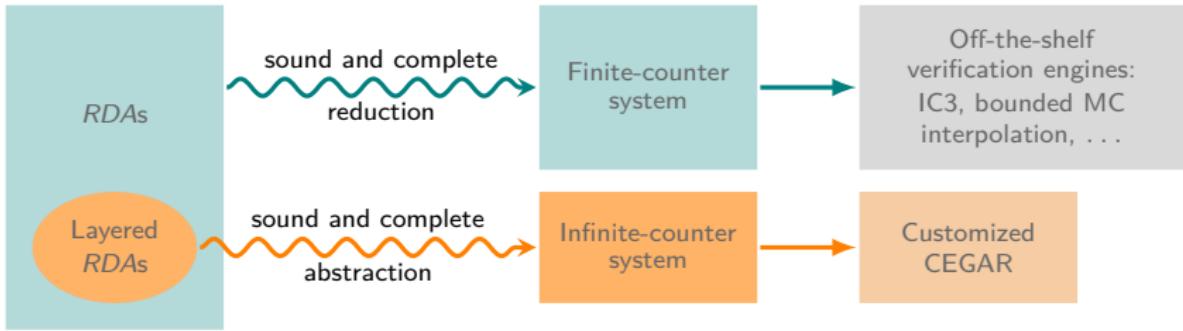
Previous work

[Bertrand, Thomas, Widder 2021] [Thomas, Sankur 2023]

# Our approach

## Challenges

- 2 sources of infinity: number of processes, number of rounds
- asynchronous communications: unbounded *drift* between processes



Previous work

[Bertrand, Thomas, Widder 2021] [Thomas, Sankur 2023]

This work

[Bertrand, Ghorpade, Rubin *under review*]

1. generalization of handled round-based distributed algorithms
2. reuse of mature model checkers e.g. nuXmv [Cavada et al. 2014]

# Outline of the rest of the talk

- 1 Modelling formalism: process template and history state-count logic
- 2 Reduction steps
- 3 Experimental validation

# Round-based process template

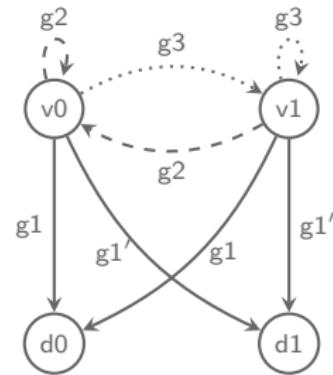
represents behaviour of a correct process

```
bool v := input_value({0,1});  
int r := 1;  
while (true) do  
    broadcast (v,r);  
    wait for n - t messages (*,r);  
    if received 2(n + t)/3 messages (w,r)  
    then d := w; halt  
    else if received (n + t)/2 messages (0,r)  
    then v := 0; r := r + 2;  
    else if received (n + t)/3 messages (1,r)  
    then v := 1; r := r + 1;  
od
```

# Round-based process template

represents behaviour of a correct process

```
bool v := input_value({0,1});  
int r := 1;  
while (true) do  
    broadcast (v,r);  
    wait for n - t messages (*,r);  
    if received 2(n + t)/3 messages (w,r)  
    then d := w; halt  
    else if received (n + t)/2 messages (0,r)  
    then v := 0; r:=r+2;  
    else if received (n + t)/3 messages (1,r)  
    then v := 1; r:=r+1;  
od
```

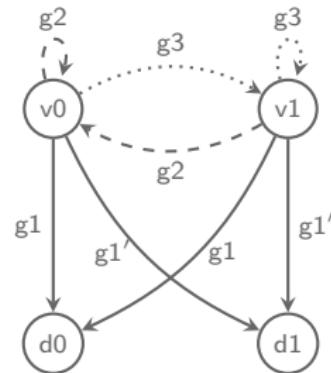


parameters  $P = \{n, t\}$   
resilience condition  $rc = n > 2t$   
locations  $\mathcal{L} = \{v0, v1, d0, d1\}$   
messages  $\mathcal{M} = \{m0, m1\}$

# Round-based process template

represents behaviour of a correct process

```
bool v := input_value({0,1});  
int r := 1;  
while (true) do  
    broadcast (v,r);  
    wait for n - t messages (*,r);  
    if received 2(n + t)/3 messages (w,r)  
    then d := w; halt  
    else if received (n + t)/2 messages (0,r)  
    then v := 0; r:=r+2;  
    else if received (n + t)/3 messages (1,r)  
    then v := 1; r:=r+1;  
od
```



parameters  $P = \{n, t\}$   
resilience condition  $rc = n > 2t$   
locations  $\mathcal{L} = \{v0, v1, d0, d1\}$   
messages  $\mathcal{M} = \{m0, m1\}$

edges w. guard and round update

broadcast associated with locations  
Bcast:  $w0 \mapsto m0$   
 $w1 \mapsto m1$   
 $d0 \mapsto \perp$   
 $d1 \mapsto \perp$

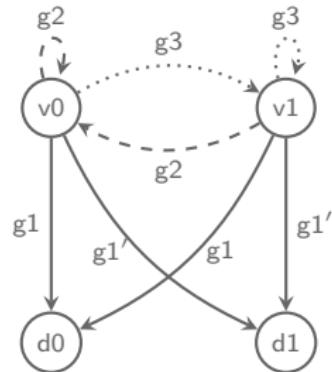
# Round-based process template

represents behaviour of a correct process

```
bool v := input_value({0,1});  
int r := 1;  
while (true) do  
    broadcast (v,r);  
    wait for n - t messages (*,r);  
    if received 2(n + t)/3 messages (w,r)  
    then d := w; halt  
    else if received (n + t)/2 messages (0,r)  
    then v := 0; r:=r+2;  
    else if received (n + t)/3 messages (1,r)  
    then v := 1; r:=r+1;  
od
```

$g1 = \text{Quorum} \wedge m0 > 2(n + t)/3$   
 $g1' = \text{Quorum} \wedge m1 > 2(n + t)/3$   
 $g2 = \text{Quorum} \wedge m0 > n + t/2$   
 $g3 = \text{Quorum} \wedge m0 \leq n + t/2 \wedge m1 \geq n + t/3$   
where Quorum =  $m0 + m1 \geq n - t$

- : no round increment
- - → : round increment of 1
- .....→ : round increment of 2



parameters  $P = \{n, t\}$   
resilience condition  $rc = n > 2t$   
locations  $\mathcal{L} = \{v0, v1, d0, d1\}$   
messages  $\mathcal{M} = \{m0, m1\}$

edges w. guard and round update

broadcast associated with locations  
Bcast:  $w0 \mapsto m0$   
 $w1 \mapsto m1$   
 $d0 \mapsto \perp$   
 $d1 \mapsto \perp$

# Semantics

## Fixed-instance semantics $\mathcal{S}(\mathcal{T}, \nu)$

- $\nu$ : fixed values for parameters ( $n$  and  $t$ )
- $n$  processes execute the same template
- configurations
  - process state: current location and round index, multiset of received messages
  - network state: multiset of broadcast messages
- actions
  - reception of a message by a process
  - process update according to a template rule (if guard permits; updates location and round index)
    - infinitely many finite-valued variables  
(counting the processes in each location and round)

Parameterized semantics  $\mathcal{S}(\mathcal{T}) = \sqcup_{\nu \models_{RC}} \mathcal{S}(\mathcal{T}, \nu)$

→ infinitely many unbounded variables

# History State-Count Logic

$$\psi ::= \forall r. \alpha_r \mid \beta \mid \neg\psi \mid \psi \wedge \psi$$

round-local atom  $\alpha_r ::= \sum_{\ell \in \mathcal{L}} c_\ell \cdot \kappa(\ell, r) \leq \varphi(n, t)$

cumulative atom  $\beta ::= \sum_{\ell \in \mathcal{X}} c_\ell \cdot \sum_{r \in \mathbb{N}} \kappa(\ell, r) \leq \varphi(n, t)$

$\varphi(n, t)$  is a linear term with variables  $n$  and  $t$

$\kappa(\ell, r)$  counts the number of process visits to location  $\ell$  in round  $r$

# History State-Count Logic

$$\psi ::= \forall r. \alpha_r \mid \beta \mid \neg\psi \mid \psi \wedge \psi$$

round-local atom  $\alpha_r ::= \sum_{\ell \in \mathcal{L}} c_\ell \cdot \kappa(\ell, r) \leq \varphi(n, t)$

cumulative atom  $\beta ::= \sum_{\ell \in \mathcal{X}} c_\ell \cdot \sum_{r \in \mathbb{N}} \kappa(\ell, r) \leq \varphi(n, t)$

$\varphi(n, t)$  is a linear term with variables  $n$  and  $t$

$\kappa(\ell, r)$  counts the number of process visits to location  $\ell$  in round  $r$

## Expressivity of HSCL

- Agreement :=  $\forall r. \kappa(d0, r) \leq 0 \vee \forall r. \kappa(d1, r) \leq 0$
- Validity :=  $\forall r. \kappa(d0, r) \leq 0$  (assuming all start with  $v = 1$ )
- Termination :=  $\neg(\sum_r \kappa(d0, r) + \kappa(d1, r) \leq N_c - 1)$
- RestrictedTermination :=  $\neg(\sum_r \kappa(d0, r) + \kappa(d1, r) \leq 0) \longrightarrow$  Term
- LeaderUniqueness :=  $\forall r. \kappa(ldr, r) \leq 1$

# Overview of reductions (1)

- Step 1: received message abstraction

- only sent messages are kept in the network state
- local counters for received messages are abstracted away
- similar in spirit to e.g. [Stoilkovska, Konnov, Widder, Zuleger 2020]
- always sound, and also complete for *common* templates  
within a round subsequent guards are monotone  
e.g.  $m_0 \geq n/3$  cannot follow  $m_0 \geq n/2$

# Overview of reductions (1)

- Step 1: received message abstraction

- only sent messages are kept in the network state
- local counters for received messages are abstracted away
- similar in spirit to e.g. [Stoilkovska, Konnov, Widder, Zuleger 2020]
- always sound, and also complete for *common* templates  
within a round subsequent guards are monotone  
e.g.  $m_0 \geq n/3$  cannot follow  $m_0 \geq n/2$

- Step 2: process identity abstraction

- process ids are irrelevant, only number of processes in each location and round matter
- classical counting abstraction from parameterized verification of systems composed of identical anonymous processes [German, Sistla 1992]

## Overview of reductions (2)

- Step 3: **synchronous restriction**

- re-ordering to focus on "semi-synchronous" executions
  - the sequence of target round indices is non-decreasing
- commutativity arguments [Chaouch-Saad, Charron-Bost, Merz 2009]

# Overview of reductions (2)

- Step 3: **synchronous restriction**

- re-ordering to focus on "semi-synchronous" executions  
the sequence of target round indices is non-decreasing
- commutativity arguments [Chaouch-Saad, Charron-Bost, Merz 2009]

- Step 4: **bounded-window abstraction**

- counters that no longer influence updates can be forgotten
- window of size  $b + 1$  suffices, with  $b$  bound on round increment in template
- sliding window can be encoded with finitely many counters

# Overview of reductions (2)

- Step 3: **synchronous restriction**

- re-ordering to focus on "semi-synchronous" executions  
the sequence of target round indices is non-decreasing
- commutativity arguments [Chaouch-Saad, Charron-Bost, Merz 2009]

- Step 4: **bounded-window abstraction**

- counters that no longer influence updates can be forgotten
- window of size  $b + 1$  suffices, with  $b$  bound on round increment in template
- sliding window can be encoded with finitely many counters

For common templates, these four steps are **sound and complete** for history state-count properties.

# From HSCL to LTL

Two more transformation steps on models: history-record extension and round identify abstraction

→ allows to reduce HSCL verification to LTL model checking

# From HSCL to LTL

Two more transformation steps on models: history-record extension and round identify abstraction

→ allows to reduce HSCL verification to LTL model checking

- Agreement :=  $\mathbf{G}(\text{local}(d0) \leq 0 \vee \text{local}(d1) \leq 0)$
- Validity :=  $\mathbf{G}(\text{local}(d0) \leq 0)$
- Termination :=  $\neg\mathbf{G}(\text{cumul}(d0) + \text{cumul}(d1) \leq N_c)$
- RestrictedTermination :=  $\mathbf{G}(\text{cumul}(d0) + \text{cumul}(d1) \leq 0)$  → Termination
- LeaderUniqueness :=  $\mathbf{G}(\text{local}(ldr) \leq 1)$

# From HSCL to LTL

Two more transformation steps on models: history-record extension and round identify abstraction

→ allows to reduce HSCL verification to LTL model checking

- Agreement :=  $\mathbf{G}(\text{local}(d0) \leq 0 \vee \text{local}(d1) \leq 0)$
- Validity :=  $\mathbf{G}(\text{local}(d0) \leq 0)$
- Termination :=  $\neg \mathbf{G}(\text{cumul}(d0) + \text{cumul}(d1) \leq N_c)$
- RestrictedTermination :=  $\mathbf{G}(\text{cumul}(d0) + \text{cumul}(d1) \leq 0)$  → Termination
- LeaderUniqueness :=  $\mathbf{G}(\text{local}(ldr) \leq 1)$

Parameterized verification of HSCL on round-based distributed algorithms reduces to LTL model checking on finite-counter systems

Rk: for fixed parameters, reduction to LTL over finite-state systems

# Experimental validation

Case studies to demonstrate applicability of the approach

- .smv file with counter system and LTL properties
- IC3 engine of nuXmv: check\_ltlspec\_ic3

3 consensus algorithms with round increment of at most 1

Protocol	loc.	rules	rc	Agree.	Valid.	Term.	R. Term.
Ben-Or (crash)	9	26	$n > 2t$	1.4s (13)	0.4s (9)	0.5 (3)	3.1s (8)
Ben-Or (Byz.)	10	27	$n > 5t$	7.0s (11)	1.2s (7)	0.6 (3)	4.3s (7)
Bracha (Byz.)	12	31	$n > 3t$	14.0s (14)	1.8s (8)	0.7 (3)	6.5s (11)

1 leader election protocol with round increment of at most 2

Protocol	b	loc.	rules	rc	Leader U.
Raft leader election	2	11	25	$n > 2t$	1.8s (8)

**Additional tests:** bugged variants (altered guards or resilience condition) detected within seconds; also verification of fixed parameter valuations (thus finite-state model checking)

# Conclusion and future work

## Contribution

- verification of correctness properties of round-based distributed algorithms
- reduction theorems to encode it into LTL verification over finite-counter system
- experiments conducted with symbolic model checker nuXmv

# Conclusion and future work

## Contribution

- verification of correctness properties of round-based distributed algorithms
- reduction theorems to encode it into LTL verification over finite-counter system
- experiments conducted with symbolic model checker nuXmv

## Remaining challenges

- how to model and verify probabilistic behaviors?
- how to deal with unbounded round jumps?
- how to deal with algorithms in which the number of locations per round grows with round index?

# Conclusion and future work

## Contribution

- verification of correctness properties of round-based distributed algorithms
- reduction theorems to encode it into LTL verification over finite-counter system
- experiments conducted with symbolic model checker nuXmv

## Remaining challenges

- how to model and verify probabilistic behaviors?
- how to deal with unbounded round jumps?
- how to deal with algorithms in which the number of locations per round grows with round index?



Vielen Dank für Ihre Aufmerksamkeit!